

Real-Software Qualification: Three Saved-State Controls

EnvLoop | v0.6.1 addendum | 24 September 2026

Study status: 0/24 researcher campaigns and 0/600 distinct official final task identities completed. The observations below are development controls, not student-model scores, hidden final tasks, or evidence of training gain.




Figure 1. Within-task positive, negative, and reset signals observed in real software. Values across rows are deliberately not pooled.

Observed readback

Surface	Evidence-bearing observation	Boundary
Microsoft Excel for the web	A five-sheet retail working-capital workbook with 112 authentic SEC companyfacts observations and 92 formula targets was edited through the GUI. Partial, full, and same-file manual reset downloads scored 1/9, 9/9, and 0/9 faults repaired. The full repair passed 92/92 targets and two evaluator-side replays.	Public development fixture. No Qwen run or automated fresh-attempt cloud reset.
Microsoft PowerPoint for the web	A 32-slide PPT-Eval selection deck was Office-normalized before scoring. Changing only one word of the title scored 1.0 upstream but 0.0 under the independent all-runs guard. The complete 48 pt title scored 1.0 under both; GUI rollback returned the target score to 0.0.	One selection task, no final task. Office's first-save visual equivalence was not comprehensively proven.
Magento admin	After repair of one isolated clone's search service and base URL, a scripted GUI control saved five green variants at \$47 from \$52. The unchanged published evaluator scored 1.0; a wrong-color save scored 0.0. Monitored SQL and all 181 search documents reset to baseline.	One clone and one public task. No model score or general 100-task reset.

What the controls establish

Excel: The operator repaired one formula, saved, reloaded, and downloaded the file; the independent oracle gave isolated 1/9 credit but failed the complete task. Continuing on the same cloud file, nine repairs passed 92 targets and both public development replay profiles. The operator then manually reverted the same file. Formula text, source values, checked structure, and tables matched the seed; all 92 Excel-populated caches were current. The reset scored 0/9. A separate full positive and automatic reset remain unverified.

PowerPoint: A frozen, evaluator-owned Office-normalized baseline prevented first-save package rewriting from being mistaken for actor damage. The unchanged PPT-Eval check nevertheless accepted a mixed-size title as a perfect edit. The independent guard required every nonempty title run to be 48 pt and rejected it. A complete GUI title selection then passed both checks, and a same-file GUI rollback restored the 42 pt title and target score of zero.

Magento: The previous task-777 rejection remains valid for its failed clone: HTTP 302 and an evaluator score of 1.0 had not persisted the prices. The repaired disposable clone produced a native save confirmation, independent SQL readback of five target prices, and a wrong-color negative. SQL rollback alone left stale search state, so the qualification tool rebuilt and compared all 181 search documents after both cases. Its nine monitored SQL business tables also matched baseline after each case.

Admission boundary

The public Excel builder exposes its fault map and replay vectors; the published PPT-Eval and WebArena task identities are also discoverable. These cases are reproducible development probes, not hidden official finals. A cell still needs source-disjoint final identities, 100 individually admitted tasks, bounded model actions, evaluator-owned fresh reset, and observed checkpoint results. No cell currently satisfies that complete gate.

Evidence ledger

The field-limited public receipts below bind saved artifact hashes, scores, and reset scope without publishing account identities, private document links, raw authenticated traces, or evaluator-only answers.

1. docs/evidence/sec-retail-working-capital-excel-web-controls-2026-09-24.json
2. docs/evidence/ppt-eval-office-web-gate-2026-09-24.json
3. docs/evidence/magento-variant-price-qualified-2026-09-24.json

This addendum supplements the v0.6 qualification note. It does not revise the immutable v0.5 Kanboard result or report the planned full-scale study.